



## GROUP-226

S.NO	NAME	BITS ID	CONTRIBUTION %
1	Suresh M	2024ac05271	100%
2	Chatrathi Lavanya	2024ac05261	100%
3	Ravi Shankar	2024ac05260	100%
4	Disha Gaikwad	2024ac05424	100%
5	Midhun S	2024ad05463	100%

```
In [1]: import numpy as np
import time

# =====#
# Configuration
# =====#

K = 3                      # Pipeline stages
D = 4                      # Data parallel replicas
M = 6                      # Micro-batches
GRAD_SIZE = 120             # Must be divisible by D
COMM_MODE = "INT8"          # "FP16" or "INT8"
LR = 0.01

# =====#
# Quantization
# =====#

def quantize_int8(t):
    max_val = np.max(np.abs(t))
    scale = max_val / 127.0 if max_val != 0 else 1.0
    q = np.round(t / scale).astype(np.int8)
    return q, scale

def dequantize_int8(q, scale):
    return q.astype(np.float32) * scale

# =====#
# Device
# =====#

class Device:
    def __init__(self, rank, stage):
        self.rank = rank
        self.stage = stage
        self.weight = np.ones(GRAD_SIZE, dtype=np.float32)
        self.local_grad = np.zeros(GRAD_SIZE, dtype=np.float32)
```

```

        self.comm_bytes = 0

# =====
# Hybrid Simulator
# =====

class HybridSimulator:

    def __init__(self):

        self.devices = []
        self.step_matrix = []

        for dp in range(D):
            for stage in range(K):
                rank = dp * K + stage
                self.devices.append(Device(rank, stage))

        self.stage_groups = {
            s: [d for d in self.devices if d.stage == s]
            for s in range(K)
        }

    # -----
    # 1F1B Pipeline Scheduling
    # -----
    def run_pipeline(self):

        total_steps = M + K - 1

        for step in range(total_steps):

            step_events = []

            for d in self.devices:

                stage = d.stage
                warmup = K - stage - 1
                event = "-"

                # Forward
                if step < M and step >= stage:
                    mb = step - stage
                    event = f"F{mb}"

                # Backward
                if step >= warmup:
                    mb = step - warmup
                    if 0 <= mb < M:
                        event = f"B{mb}"
                        grad = np.random.randn(GRAD_SIZE).astype(np.float32)
                        d.local_grad += grad

```

```

        step_events.append((d.rank, event))

    self.step_matrix.append(step_events)

# -----
# Chunked Ring AllReduce
# -----
def ring_all_reduce(self, stage):

    group = self.stage_groups[stage]
    grads = np.stack([d.local_grad for d in group])
    chunks = np.split(grads, D, axis=1)

    # Reduce-Scatter
    for step in range(D - 1):
        for i in range(D):
            src = i
            dst = (i + 1) % D
            chunk_idx = (i - step) % D

            send_chunk = chunks[chunk_idx][src]

            if COMM_MODE == "FP16":
                payload = send_chunk.astype(np.float16)
                group[src].comm_bytes += payload.nbytes
                recv_chunk = payload.astype(np.float32)
            else:
                q, scale = quantize_int8(send_chunk)
                group[src].comm_bytes += q.nbytes
                recv_chunk = dequantize_int8(q, scale)

            target_idx = (dst - step - 1) % D
            chunks[target_idx][dst] += recv_chunk

    # All-Gather
    for step in range(D - 1):
        for i in range(D):
            src = i
            dst = (i + 1) % D
            chunk_idx = (i - step) % D

            send_chunk = chunks[chunk_idx][src]

            if COMM_MODE == "FP16":
                payload = send_chunk.astype(np.float16)
                group[src].comm_bytes += payload.nbytes
                recv_chunk = payload.astype(np.float32)
            else:
                q, scale = quantize_int8(send_chunk)
                group[src].comm_bytes += q.nbytes
                recv_chunk = dequantize_int8(q, scale)

```

```

        target_idx = (dst - step - 1) % D
        chunks[target_idx][dst] = recv_chunk

    final = np.concatenate(chunks, axis=1) / D

    for i, d in enumerate(group):
        d.local_grad = final[i]

# -----
def optimizer_step(self):

    for stage in range(K):
        for d in self.stage_groups[stage]:
            d.weight -= LR * d.local_grad
            d.local_grad[:] = 0

# -----
def evaluate_accuracy(self):

    np.random.seed(42)
    X = np.random.randn(300, GRAD_SIZE).astype(np.float32)
    true_w = np.random.randn(GRAD_SIZE).astype(np.float32)

    logits = X @ true_w
    y_true = (logits > 0).astype(int)

    model_w = self.stage_groups[0][0].weight
    preds = X @ model_w
    y_pred = (preds > 0).astype(int)

    return np.mean(y_pred == y_true) * 100

# -----
def total_comm_bytes(self):
    return sum(d.comm_bytes for d in self.devices)

# -----
def print_generation_timeline(self):

    print("\n===== PIPELINE GENERATION TIMELINE =====")

    header = "Step | " + " | ".join(
        [f"R{d.rank}(S{d.stage})" for d in self.devices]
    )
    print(header)
    print("-" * len(header))

    for step, events in enumerate(self.step_matrix):
        row = f"{step:<4} | "
        for _, event in events:
            row += f"{event:^10}| "
        print(row)

```

```

print("\n=====\n")

# -----
def verify(self):

    print("Replica Consistency Check:")
    for stage in range(K):
        weights = [d.weight for d in self.stage_groups[stage]]
        diffs = [np.linalg.norm(weights[0] - w) for w in weights]
        print(f" Stage {stage}: Max Weight Difference = {max(diffs)}")

# -----
def run(self):

    print(f"\nHybrid Parallel Simulation (K={K}, D={D})")
    print(f"Total Nodes = {K * D}")
    print(f"Communication Mode = {COMM_MODE}\n")

    start_time = time.perf_counter()

    self.run_pipeline()

    for stage in range(K):
        self.ring_all_reduce(stage)

    self.optimizer_step()

    end_time = time.perf_counter()
    runtime = end_time - start_time

    accuracy = self.evaluate_accuracy()

    print("\n===== PERFORMANCE REPORT =====\n")
    print(f"Runtime (seconds): {runtime:.6f}")
    print(f"Prediction Accuracy: {accuracy:.2f}%")
    print(f"Total Communication Bytes: {self.total_comm_bytes()}")
    print("\n======\n")

    self.verify()
    self.print_generation_timeline()

# =====
if __name__ == "__main__":
    HybridSimulator().run()

```

Hybrid Parallel Simulation (K=3, D=4)  
Total Nodes = 12  
Communication Mode = INT8

===== PERFORMANCE REPORT =====

Runtime (seconds): 0.016619  
Prediction Accuracy: 49.00%  
Total Communication Bytes: 2160

=====

Replica Consistency Check:

Stage 0: Max Weight Difference = 0.0003009646898135543  
Stage 1: Max Weight Difference = 0.0002646037610247731  
Stage 2: Max Weight Difference = 0.0002457082155160606

===== PIPELINE GENERATION TIMELINE =====

Step	R0(S0)	R1(S1)	R2(S2)	R3(S0)	R4(S1)	R5(S2)	R6(S0)	R7(S1)	R8(S2)	R9(S0)	R10(S1)	R11(S2)
------	--------	--------	--------	--------	--------	--------	--------	--------	--------	--------	---------	---------

-----													
0		F0		-		B0		F0		-		B0	
F0		-		B0		F0		-		B0			
1		F1		B0		B1		F1		B0		B1	
F1		B0		B1		F1		B0		B1			
2		B0		B1		B2		B0		B1		B2	
B0		B1		B2		B0		B1		B2			
3		B1		B2		B3		B1		B2		B3	
B1		B2		B3		B1		B2		B3			
4		B2		B3		B4		B2		B3		B4	
B2		B3		B4		B2		B3		B4			
5		B3		B4		B5		B3		B4		B5	
B3		B4		B5		B3		B4		B5			
6		B4		B5		-		B4		B5		-	
B4		B5		-		B4		B5		-			
7		B5		-		-		B5		-		-	
B5		-		-		B5		-		-			

=====

In [ ]: